

Introduction to Kernel Programming

Luca Abeni

`luca.abeni@santannapisa.it`

About the Course

- Goal: understand how to code an OS kernel
 - This course will introduce the students to the pains and joys of kernel programming...
 - Of course, this is only an introduction!
- First question: so, why is kernel programming so different?
 - But, first of all, what is a kernel?
 - ...And what is an Operating System?
 - ...And, what is a computer? Where are we coming from? Where are we going?
 - No, I will not answer the last questions...

Practical Details

- The course is 20 hours long (2 credits)
- Organized in lessons by 2:30 hours each
- Some theory and some practice
 - There cannot be practice without theory...
 - ...But theory is useless without practice!
- For informations, email
`luca.abeni@santannapisa.it`
 - Office: TeCIP (via Moruzzi 1)

Overview

- Quick recap about architecture, OSs, kernels
 - Privileged instructions → kernel
- The kernel execution environment
- Kernel development: DIY, or existing systems (Linux)?
- Introduction to the Linux kernel
 - Kernel modules, concurrency, synchronization, ...
- Some examples

The Operating System

- Operating System: set of computer programs acting as an interface between applications and hardware
 - “Set of computer programs”: the OS is not a single program!
 - “Acting as an interface... ”: applications do not directly access the hardware (must use the OS)
- So, the OS:
 - Hides the hardware details to user applications
 - Controls the execution of user programs
 - Manages the hardware and software resources
- Applications running on an OS can use the CPU Assembly extended with some additional instructions: the *system calls*

The OS Kernel

- Modern CPUs: different privilege levels (user level and privileged level)
 - Actually, it can be even more complex than this
← hypervisor mode, ...
- Security / protection → only a small amount of **trusted** code should run with a high privilege level
 - OS Kernel: part of the OS executing with a high privilege level
- Regular user applications execute at a lower privilege level
 - To protect the system from malicious programs

The OS Kernel - Again

- Kernel: part of the OS running at high privilege level
 - Can do (almost!) everything: even crash the system
- This is why it must be trusted... Very critical component of the system
 - Security and stability depend on it!
 - But also the system performance depends on it...
- “With Great Power Comes Great Responsibility”
- Applications rely on the kernel to do everything important / critical
 - How is kernel execution invoked? Interrupts / exceptions (hw or sw)

Kernel Responsibilities

- Multiprogramming:
 - Multiple tasks (processes or threads) on few CPUs
 - Memory protection: multiple address spaces (paging, segmentation)
 - CPU privilege levels: system and user
- Low level hardware details:
 - Interrupt handling, boot, device drivers, system calls, ...
 - Important data structures (memory page tables, ...)
- Kernel address space: can see all the system memory

Kernel Functionalities

- System boot → configure and set-up the system so that virtual memory and multitasking can work
 - Configure memory, page tables & friends
 - Once it is done, start the first user-space process: in Unix, it is `init` (PID 1)
 - Then, execution returns to the kernel only through interrupts
- Hardware interrupt handlers (ISRs, used by drivers)
- System calls (software interrupts, ...)
 - Interrupts cause a privilege change (user → system)
 - Syscalls and ISRs can cause context switches

Task Handling

- The kernel handles processes and threads
- Each task is characterised by:
 - A Task Descriptor (TSS)
 - A Task Body (code implementing the task)
 - Some (public or private) data (Address Space)
- Task Descriptor → contains copy of the CPU registers, including:
 - Pointers to user-space and kernel stack
 - Address Space configuration (CR3, ...)
 - Protection level (CPL)
- The task body is technically part of the address space

Task Address Space

- Divided in **segments**:
 - Code Segment (task body)
 - Data Segment
 - Initialised Data, BSS, Heap
 - Stack Segment
- Recently, some additional segments (RO data, etc...)
- Before starting a task, the OS kernel has to:
 - Initialise memory segments (allocate virtual memory)
 - Allocate and set up a stack, initialise the stack pointer, etc...
 - Allocate and initialise (to 0) the BSS

The Rest of the OS

- The kernel is only part of the OS
- There also are many user-space components
 - System libraries
 - System programs
 - ...
- System libraries → needed to properly invoke kernel functionalities (hide the syscall mechanism)
- System programs → needed to properly boot and use the system

How does a Kernel Look Like

- No single entry point
 - “Boot entry point” + system calls
- Kernel Memory Address Space: all the memory can be accessed
 - No memory protection!!!
 - Kernel-space code can easily corrupt important data structures!
- No standard runtime
 - C code cannot include `<stdio.h>` and friends...
 - No standard C library!

Kernel Programming

- The kernel must provide the utility functions to be used
 - Example, no `printf()`, but `printk()`...
- Errors do not result in segmentation faults...
- ...But can cause system crashes!
- Other weird details
 - No floating point (do not use `float` or `double`)
 - Small stack (*4KB* or *8KB*)
 - Atomic contexts, ...

Kernel Programming Language

- OS kernels can be coded in many different languages
 - But some amount of Assembly is needed...
- For some languages, additional restrictions apply (example: for C++, generally no RTTI)
 - And in order to use some languages the kernel must implement a large runtime...
- Kernels are generally coded in C or C++ (with restrictions)
 - After all, the C language has been invented exactly for this purpose!
- Simplest choice: C + some Assembly (inline and not)

Example: the Linux Kernel

- The Linux kernel uses C
- Subset of C99 + some extensions (`likely()` / `unlikely()` annotations, etc...)
- As said, no access to standard libraries
 - Different set of header files and utility functions
- Strict coding style to control the quality (`Documentation/CodingStyle`, `scripts/checkpatch.pl`)
- Some Assembly is used (for entry points, etc...)
- Example: Linked Lists (`include/linux/list.h`)

First Adventures in Kernel Land

- First experiments with kernel programming
 - Should we write our own kernel?
 - Or use an existing kernel as a basis?
- Our own didactic kernel: simpler, maybe we learn more...
 - ...But we can easily get lost in low-level hw details!!!
- Work on an existing kernel: it might be more complex...
 - ...But we can focus on the aspects we are interested in
 - The rest of the kernel already exists!